

Data Mining and Machine learning

Naïve Bayes classifier

Edgar Acuna

Department of Mathematical Science

academic.uprm.edu/eacuna

How to build a Bayes Classifier

- Assume you want to predict output Y which has arity n_Y and values $V_1, V_2, \dots V_{n_Y}$.
- Assume there are m input attributes called $X_1, X_2, \dots X_m$
- Break dataset into n_Y smaller datasets called $DS_1, DS_2, \dots DS_{n_Y}$.
- Define DS_i = Records in which $Y=v_i$
- For each DS_i , learn Density Estimator M_i to model the input distribution among the $Y=v_i$ records.

How to build a Bayes Classifier

- Assume you want to predict output Y which has arity n_Y and values $V_1, V_2, \dots V_{n_Y}$.
- Assume there are m input attributes called $X_1, X_2, \dots X_m$
- Break dataset into n_Y smaller datasets called $DS_1, DS_2, \dots DS_{n_Y}$.
- Define DS_i = Records in which $Y=v_i$
- For each DS_i , learn Density Estimator M_i to model the input distribution among the $Y=v_i$ records.
- M_i estimates $P(X_1, X_2, \dots X_m \mid Y=v_i)$

How to build a Bayes Classifier

- Assume you want to predict output Y which has arity n_Y and values V_1, V_2, \dots, V_{n_Y} .
 - Assume there are m input attributes called X_1, X_2, \dots, X_m
 - Break dataset into n_Y smaller datasets called $DS_1, DS_2, \dots, DS_{n_Y}$.
 - Define DS_i = Records in which $Y=v_i$
 - For each DS_i , learn Density Estimator M_i to model the input distribution among the $Y=v_i$ records.
 - M_i estimates $P(X_1, X_2, \dots, X_m \mid Y=v_i)$
- Idea: When a new set of input values ($X_1 = u_1, X_2 = u_2, \dots, X_m = u_m$) come along to be evaluated predict the value of Y that makes $P(X_1, X_2, \dots, X_m \mid Y=v_i)$ most likely

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(X_1 = u_1 \cdots X_m = u_m \mid Y = v)$$

Is this a good idea?

How to build a Bayes Classifier

- Assume you want to predict output Y which has arity n_Y and values V_1, V_2, \dots, V_{n_Y} .
 - Assume there are m input attributes called X_1, X_2, \dots, X_m .
 - Break dataset into n_Y smaller datasets called DS_i .
 - Define DS_i = Records in which $Y=v_i$.
 - For each DS_i , learn Density Estimator M_i for the joint distribution among the $Y=v_i$ records.
 - M_i estimates $P(X_1, X_2, \dots, X_m \mid Y=v_i)$.
- Idea: When a new set of input values $(X_1 = u_1, X_2 = u_2, \dots, X_m = u_m)$ come along to be evaluated, predict the value of Y that makes $P(Y=v_i \mid X_1, X_2, \dots, X_m)$ most likely

Much Better Idea

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(Y = v \mid X_1 = u_1 \cdots X_m = u_m)$$

Is **this** a good idea?

Bayes Classifiers

1. Learn the distribution over inputs for each value Y .
2. This gives $P(X_1, X_2, \dots, X_m \mid Y=v_i)$.
3. Estimate $P(Y=v_i)$ as fraction of records with $Y=v_i$.
4. For a new prediction:

$$\begin{aligned} Y^{\text{predict}} &= \operatorname{argmax}_v P(Y = v \mid X_1 = u_1 \cdots X_m = u_m) \\ &= \operatorname{argmax}_v P(X_1 = u_1 \cdots X_m = u_m \mid Y = v) P(Y = v) \end{aligned}$$

Bayes Classifiers

1. Learn the distribution over inputs for each value Y .

2. This gives $P(X_1, X_2, \dots, X_m / Y=v_i)$.

3. Estimate $P(Y=v_i)$ as fraction of records with $Y=v_i$.

4. For a new prediction:

$$Y^{\text{predict}} = \operatorname{argmax}_v P(Y = v \mid X_1 = u_1, \dots, X_m = u_m) \\ = \operatorname{argmax}_v P(X_1 = u_1 \cdots X_m = u_m \mid Y = v)$$

We can use our favorite
Density Estimator here.
But the easier and faster is
•Naïve Density Estimator

Naïve Bayes Classifier

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(X_1 = u_1 \cdots X_m = u_m \mid Y = v)P(Y = v)$$

In the case of the naive Bayes Classifier this can be simplified:


$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(Y = v) \prod_{j=1}^{n_Y} P(X_j = u_j \mid Y = v)$$

In here, the random variables X_1, \dots, X_m are considered independent

Naïve Bayes Classifier

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(X_1 = u_1 \cdots X_m = u_m \mid Y = v) P(Y = v)$$

In the case of the naive Bayes Classifier this can be simplified:

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(Y = v) \prod_{j=1}^{n_Y} P(X_j = u_j \mid Y = v)$$


If you have a lot of input attributes **that** product will underflow in floating point math. You should use logs:

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} \left(\log P(Y = v) + \sum_{j=1}^{n_Y} \log P(X_j = u_j \mid Y = v) \right)$$

Naïve Bayes classifier

Naïve Bayes classifiers can be applied when there are continuous attributes in the dataset. There are two options:

- a) Discretize the continuous attributes using methods such as, Equal width discretization, Equal frequency discretization and, Entropy with MDLP discretization. Scikit learn has two functions BernoulliNB and MultinomialNB.
- b) Assume a Gaussian distribution for the continuous predictors. Scikit learn has a GaussianNB function.

The option a) is preferred by many.

Example

X1	X2	X3	Y
0	0	1	0
0	1	0	0
1	1	0	0
0	0	1	1
1	1	1	1
0	0	1	1
1	1	0	1

Example: Cont

To which class will be assigned the record (0,0,1)?

$$P(Y = 0) = 3/7$$

$$P(Y = 1) = 4/7$$

$$\begin{aligned} P(X_1 = 0, X_2 = 0, X_3 = 1 / Y = 0) &= P(X_1 = 0 / Y = 0) P(X_2 = 0 / Y = 0) P(X_3 = 1 / Y = 0) = \\ &= (2/3)(1/3)(1/3) = 2/27 \end{aligned}$$

$$\begin{aligned} P(X_1 = 0, X_2 = 0, X_3 = 1 / Y = 1) &= P(X_1 = 0 / Y = 1) P(X_2 = 0 / Y = 1) P(X_3 = 1 / Y = 1) = \\ &= (2/4)(2/4)(3/4) = 3/16 \end{aligned}$$

X1=0, X2=0, x3=1 will be assigned to class 1

Example 2. (continuous and discrete attributes)

X1	X2	X3	X4	Y
0	0	1	3.15	0
0	1	0	8.17	0
1	1	0	5.72	0
0	0	1	7.16	1
1	1	1	9.32	1
0	0	1	12.81	1
1	1	0	15.48	1

Example 2. (cont.)

➤ #Metodo 1. Discretizando la columna 4

> dnaiveeje2

	col1	col2	col3	col4	col5
[1,]	0	0	1	1	0
[2,]	0	1	0	1	0
[3,]	1	1	0	1	0
[4,]	0	0	1	1	1
[5,]	1	1	1	2	1
[6,]	0	0	1	2	1
[7,]	1	1	0	2	1

➤ #Metodo 2. Sin discretizar la columna 4

➤ Media y desviacion estandar de la col4 en cada clase

> mean(naiveeje2[naiveeje2[,5]==0,4])

➤ [1] 5.68

➤ > mean(naiveeje2[naiveeje2[,5]==1,4])

➤ [1] 11.1925

Example 2. (cont.)

➤ > sd(naiveeje2[naiveeje2[,5]==0,4])

➤ [1] 2.510239

➤ > sd(naiveeje2[naiveeje2[,5]==1,4])

➤ [1] 3.686293

> # a que clase sera asignado el vector(0,0,1,4.25)?

Hay que calcular $P[X1=0/Y=0]P[X2=0/Y=0]P[x3=1/Y=0]f(x4=4.25/Y=0)[Y=0]$
y compararla con

$P[X1=0/Y=1]P[X2=0/Y=1]P[x3=1/Y=1]f(x4=4.25/Y=1)[Y=1]$

> (2/27)*dnorm(4.25,5.68,2.5102)*3/7

[1] 0.009030122

> (3/16)*dnorm(4.25,11.1925,3.6862)*4/7

[1] 0.003195506

Luego el vector sera asignado a la clase 0.

Example 2. (cont.)

```
> pred=predict(a,naivebayes21[,-5],type="raw")
➤ pred1=max.col(pred)
➤ > pred1
➤ [1] 1 1 1 2 2 2 2
> table(pred1,naivebayes21[,5])
Pred1  0 1
      1 3 0
      2 0 4
Error=0
```


Naïve Bayes for Diabetes

Without Discretization

```
# Calculo de las probabilidades posteriores y de las clases predichas
clf = GaussianNB()
clf.fit(X1,y1)
proba=pd.DataFrame(clf.predict_proba(X1))
pred=clf.predict(X1)
print 'Accuracy by Resubstitution',(pred==y1).sum()/float(768)
Accuracy by Resubstitution 0.76171875
# Calculo de las probabilidades posteriores y de las clases predichas
```

With Discretization; Equal width

```
from sklearn.preprocessing import KBinsDiscretizer
est = KBinsDiscretizer(n_bins=10, encode='ordinal', strategy='uniform')
est.fit(X)
Xt = est.transform(X)
clf=MultinomialNB()
clf.fit(Xt,y1)
scores = cross_val_score(clf, X1, y1, cv=10)
print("Accuracy by cross validation: %0.3f (+/- %0.3f)" % (scores.mean(), scores.std()))
Accuracy by cross validation: 0.654 (+/- 0.048)
```

The Auto-mpg dataset

Donor: Quinlan, R. (1993)

Number of Instances: 398 minus 6 missing=392 (training: 196 test: 196):

Number of Attributes: 9 including the class attribute 7. Attribute Information:

1. mpg: continuous (discretizado $\text{bad} \leq 25, \text{good} > 25$)
2. cylinders: multi-valued discrete
3. displacement: continuous (discretizado $\text{low} \leq 200, \text{high} > 200$)
4. horsepower: continuous ((discretizado $\text{low} \leq 90, \text{high} > 90$)
5. weight: continuous (discretizado $\text{low} \leq 3000, \text{high} > 3000$)
6. acceleration: continuous (discretizado $\text{low} \leq 15, \text{high} > 15$)
7. model year: multi-valued discrete (discretizado 70-74, 75-77, 78-82)
8. origin: multi-valued discrete
9. car name: string (unique for each instance)

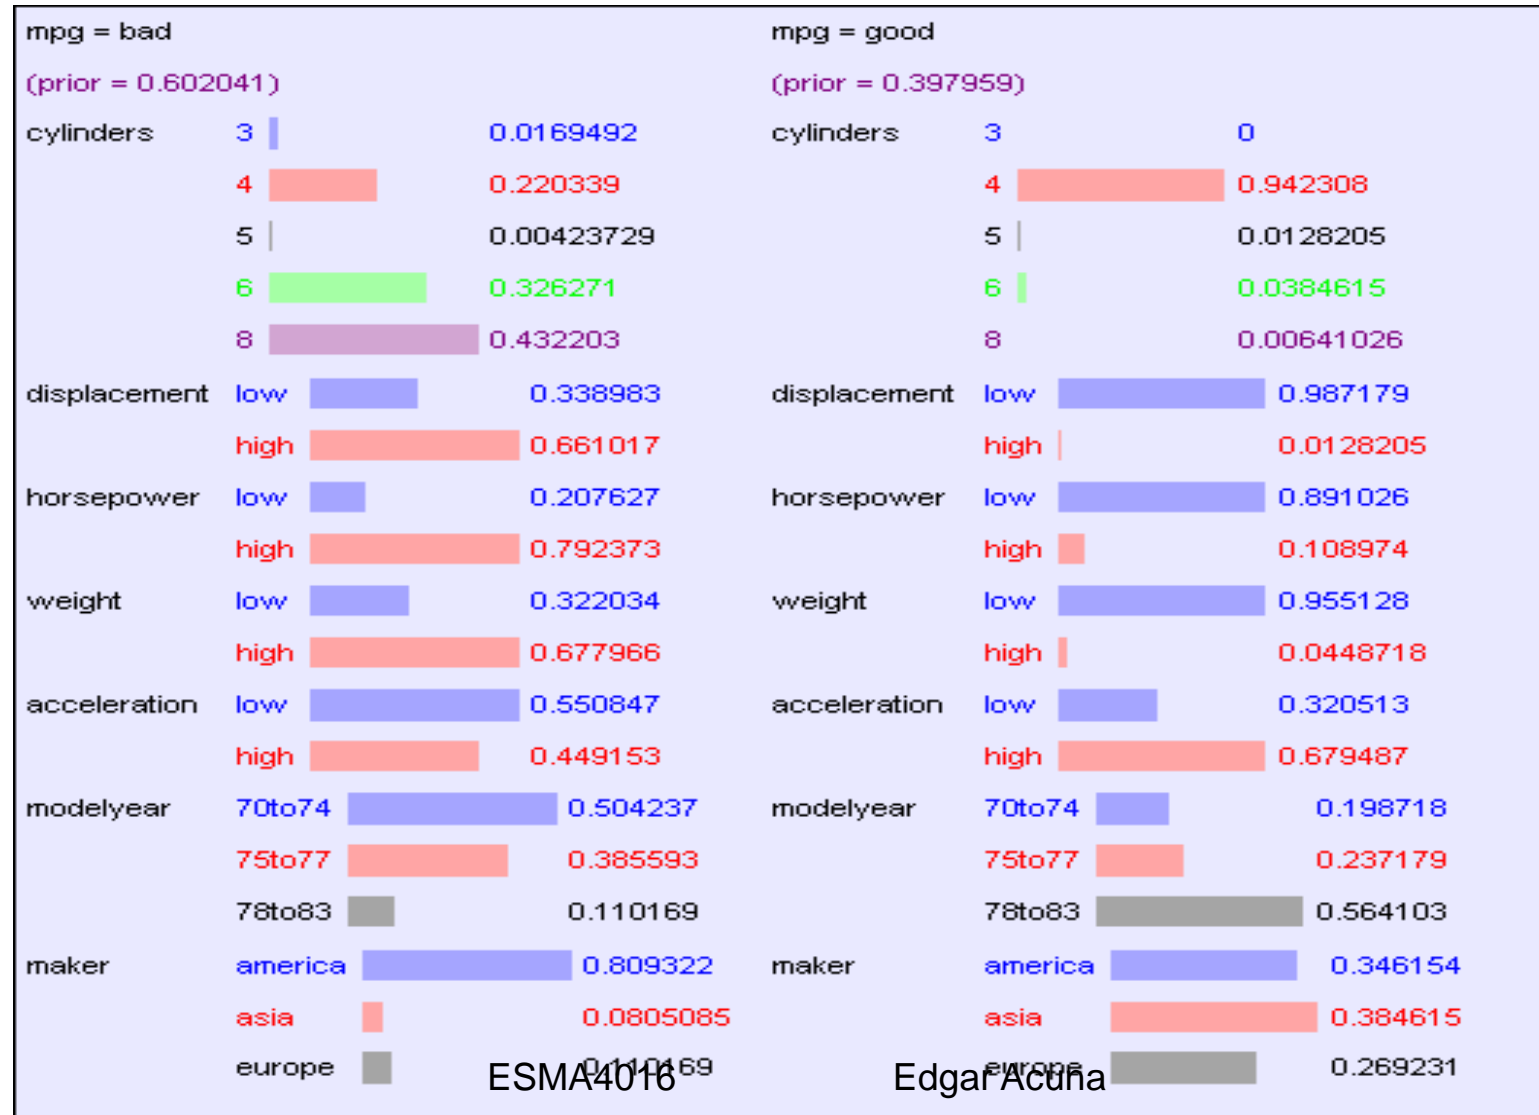
Note: horsepower has 6 missing values

Available at: <https://archive.ics.uci.edu/ml/machine-learning-databases/auto-mpg/auto-mpg.data>

The auto-mpg dataset

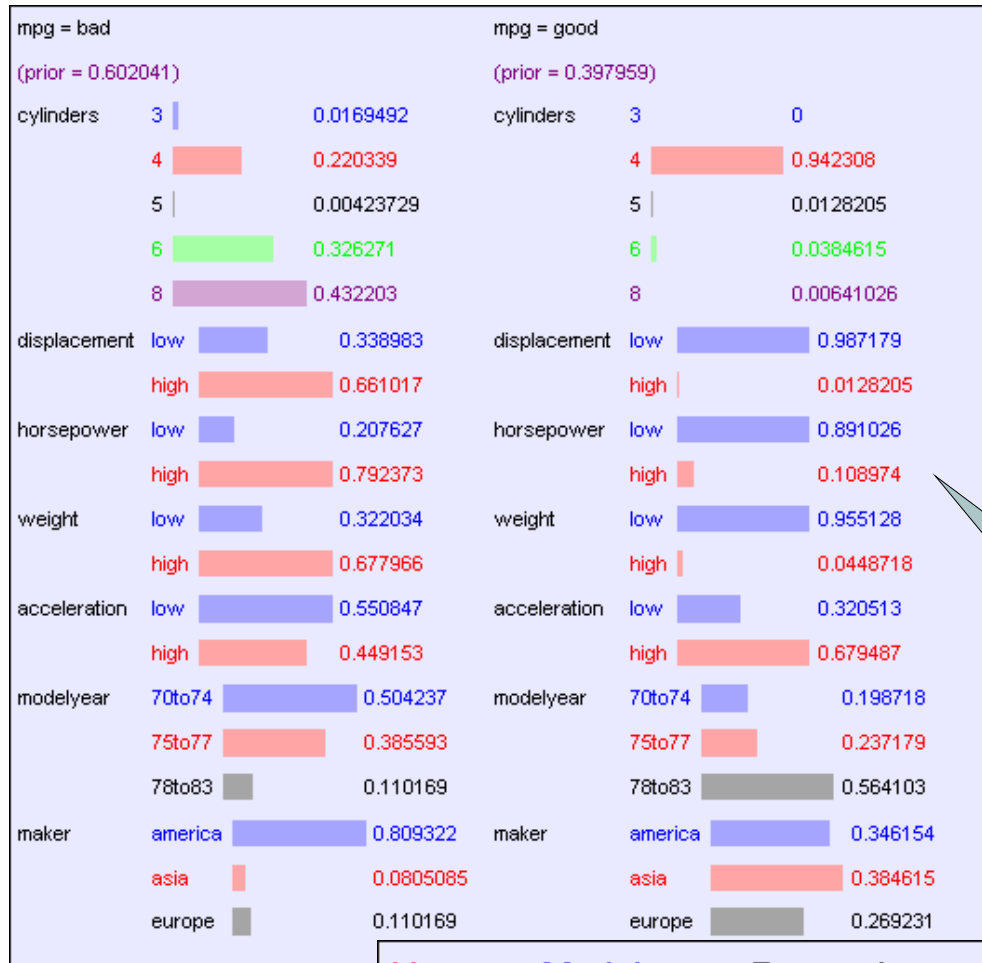
```
18.0 8 307.0 130.0 3504. 12.0 70 1 "chevrolet chevelle malibu"  
15.0 8 350.0 165.0 3693. 11.5 70 1 "buick skylark 320"  
18.0 8 318.0 150.0 3436. 11.0 70 1 "plymouth satellite"  
16.0 8 304.0 150.0 3433. 12.0 70 1 "amc rebel sst"  
17.0 8 302.0 140.0 3449. 10.5 70 1 "ford torino"  
  
.....  
27.0 4 140.0 86.00 2790. 15.6 82 1 "ford mustang gl"  
44.0 4 97.00 52.00 2130. 24.6 82 2 "vw pickup"  
32.0 4 135.0 84.00 2295. 11.6 82 1 "dodge rampage"  
28.0 4 120.0 79.00 2625. 18.6 82 1 "ford ranger"  
31.0 4 119.0 82.00 2720. 19.4 82 1 "chevy s-10"
```

Resultados del clasificador NB para “MPG”: 392 records



NB results for “mpg”

```
➤ #without discretization:
➤ > b=naiveBayes(mpg~.,data=autompg)
> pred=predict(b,autompg[,-1],type="raw")
> pred1=max.col(pred)
> table(pred1,autompg[,1])
pred1  1  2
  1 180  8
  2  56 148
> error=64/392
> error
[1] 0.1632
> #with manual discretization
> b=naiveBayes(mpg~.,data=autompg2)
> pred=predict(b,autompg2[,-1])
> table(pred,autompg2[,1])
pred  1  2
  1 182  7
  2  54 149
> 61/392
[1] 0.1556122
```



BC Results:

“MPG”: 392 records

The Classifier learned by “Naive BC”

Name	Model	Parameters	FracRight	
Model1	bayesclass	density=joint submodel=gauss gausstype=general	0.885256 +/- 0.0247796	
Model2	bayesclass	density=naive submodel=gauss gausstype=general	0.852372 +/- 0.0400495	

Advantages/Disadvantages of NB

Probabilities zeros can affect the NB classifier. A correction factor is apply to to the computation of probabilities to solve this problem.

- The discretization process seems to affect the performance of the classifier.
- Naïve Bayes is very cheap. It does not have problem working with 10,000 attributes.
- Naïve Bayes is a particular case of Bayesian networks